

POP77014: Assignment 2

Imelda Finn (22334657)

08/04/2023

Overview

This report analyses the evolution of all cause mortality in the USA 2015-2023.

1. The data is from the OECD website: COVID-19 Health Indicators, Mortality (by week)(OECD 2023)

Mortality rates have been generally decreasing over time, i.e. life expectancy has been rising. (This may not continue in the future, particularly in developed nations, as lifestyle factors such as obesity may tend to reduce inter-generational life expectancy.)

In a stable population (where births/migration replace deaths), mortality rates will be gradually reducing, as life expectancy causes fewer deaths at each age. In an expanding population, average age is decreasing so mortality rates reduce faster. In a contracting population, the average age is increasing as deaths are not balanced by births or inward (younger) migrants, so mortality rates will be steady or even rising (temporarily).

COVID-19 started in 2019, the first deaths in America were in 2020; this will distort the mortality predictions for 2021 and beyond. Some models may be better at capturing what happened, and so might be more useful in future pandemics.

Final Model: ensemble - mean of 7 models

- HoltWinters with high α ($=0.94$), quickly takes surge in mortality in 2020 into account

“if you want to make prediction intervals for forecasts made using exponential smoothing methods, the prediction intervals require that the forecast errors are uncorrelated and are normally distributed with mean zero and constant variance.” (Coghlan 2023) Assumption doesn’t hold with this data

3. If relevant, all estimated equations associated with constructing forecasts from this method
4. Report the MAPE and MAE for the training period and the validation period. You may also report other metrics if relevant.
6. A single figure showing the fit of the final version of the model to the entire period available in the data (i.e., in-sample fit. For options 1 and 2, you do not have access to the “future” data).

```
# read in the data
# show_col_types = FALSE


#https://stats.oecd.org/Index.aspx?DataSetCode=HEALTH_MORTALITY
mort <- read_csv("data/HEALTH_MORTALITY_all.csv")

## Rows: 268379 Columns: 15
## -- Column specification -----
## Delimiter: ","
## chr (10): COUNTRY, Country, GENDER, Gender, AGE, Age, VARIABLE, Variable, Fl...
## dbl (5): WEEK, Week number, YEAR, Year, Value
##
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
mort$date <- ISOweek2date(paste0(mort$YEAR, "-W", sprintf("%02d", mort$WEEK), "-4"))
mort <- mort %>%
  filter(Age == "Total" & Gender == "Total" & VARIABLE == "ALLCAUNB") %>%
  arrange(desc(COUNTRY), date)

#summary(mort)
#unique(mort$COUNTRY)

mort %>% ggplot(aes(x=date, y=Value, colour = COUNTRY)) + geom_line() +
  ylab("Deaths (all cause)")
```



```
country_code <- "USA"
country <- mort %>% filter(COUNTRY == country_code) %>% select(date, Value, YEAR)
country %>% arrange(date)
```

```
## # A tibble: 425 x 3
##   date      Value  YEAR
##   <date>    <dbl> <dbl>
## 1 2015-01-01 61763 2015
## 2 2015-01-08 61163 2015
## 3 2015-01-15 58652 2015
## 4 2015-01-22 57297 2015
## 5 2015-01-29 57367 2015
## 6 2015-02-05 56668 2015
## 7 2015-02-12 56678 2015
```

```

## 8 2015-02-19 56334 2015
## 9 2015-02-26 56509 2015
## 10 2015-03-05 55741 2015
## # i 415 more rows

length(country$date)

## [1] 425

min(country$date)

## [1] "2015-01-01"

# first week in data is 1, 2015
# 2020 has 53 weeks

# set the parameters for the time series
startDate <- min(country$date)
startYear <- year(startDate)
startMonth <- month(startDate)
startWeek <- week(startDate)
startDay <- day(startDate)

endDate <- max(country$date)
endYear <- year(endDate)
endMonth <- month(endDate)
endWeek <- week(endDate)

# specify the forecasting parameters
# solve for recommended
# look for smaller alphas to smooth out effect of pandemic
ALPHA <- 0.95
FREQ<- 52
WEEKS <- length(country$Value)

future <- as.integer(WEEKS*0.2) # 85 for full data

fivenum(country$Value)

## [1] 48194 51838 56680 60260 87415

# default graph labels
mtitle <- paste0(country_code," all causes deaths (weekly)")
stitle <- paste0(startMonth, "/", startYear, " - ", endMonth, "/",endYear)
y_lab <- "Deaths"
x_lab <- "Year"
us_y_lim <- c(48000, 90000) # check when incorporate leap year
y_lim <- us_y_lim
x_lim <- c(startDate, endDate)

Data is weekly, 425 weeks from 2015-01-01 to 2023-02-23 inclusive.

#Convert the death numbers to a time series
country.ts <- ts(country$Value, start=c(startYear,startWeek), frequency=FREQ)

print(tsfeatures(country.ts))

## # A tibble: 1 x 20

```

```
## frequency nperiods seasonal_period trend spike linearity curvature e_acf1
## <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 52 1 52 0.740 0.00000121 12.1 0.226 0.940
## # i 12 more variables: e_acf10 <dbl>, seasonal_strength <dbl>, peak <dbl>,
## # trough <dbl>, entropy <dbl>, x_acf1 <dbl>, x_acf10 <dbl>, diff1_acf1 <dbl>,
## # diff1_acf10 <dbl>, diff2_acf1 <dbl>, diff2_acf10 <dbl>, seas_acf1 <dbl>
```

```
summary(country.ts)
```

```
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 48194 51838 56680 57863 60260 87415
```

```
head(tail(country,85),1) #2021-07-15 # first week of test set
```


```
## # A tibble: 1 x 3
## date Value YEAR
## <date> <dbl> <dbl>
## 1 2021-07-15 57253 2021
```

```
# plot the whole time series
```

```
autoplot(country.ts) +
  ggtitle(mtitle, subtitle = stitle) +
  xlab(x_lab) +
  ylab(y_lab)
```

USA all causes deaths (weekly)

1/2015 – 2/2023



```
nValid <- future
nTrain <- length(country.ts) - nValid
train.ts <- window(country.ts, start = c(startYear, startWeek),
```

```

end = c(startYear, nTrain))
valid.ts <- window(country.ts, start = c(startYear, nTrain+1),
end = c(startYear, nTrain+nValid))

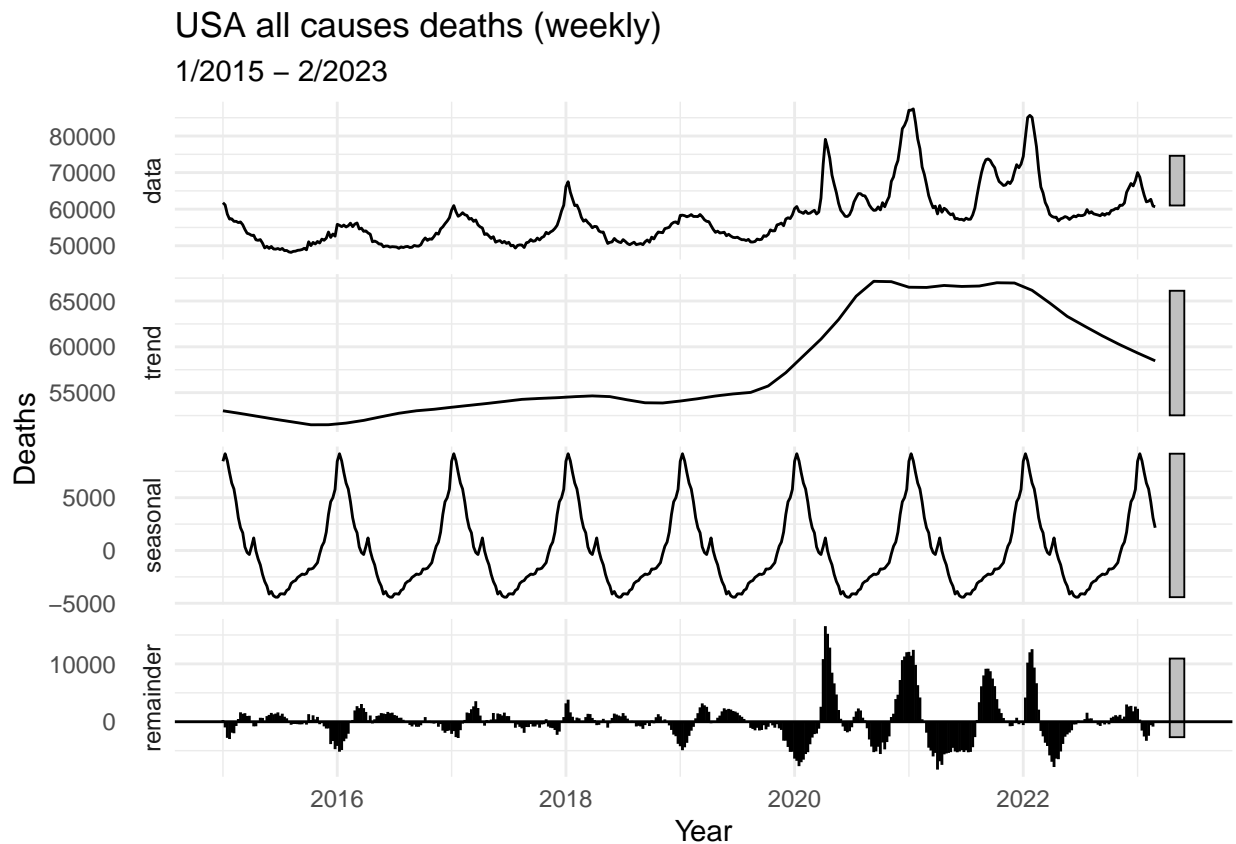
```

Consider the timeseries decomposition: seasonal component; the trend; and the remainder.

```

#plot the decomposition
(country.stl <- country.ts %>%
  stl(s.window="periodic"))%>%
  autoplot() +
  ggtitle(mtitle, subtitle = stitle) +
  xlab(x_lab) +
  ylab(y_lab)

```



value of time series at time $t = y_t = T_t + S_t + R_t$

```

# This toggles plots from plotly (interactive) to ggplot (static)
interactive <- FALSE

#country %>% plot_time_series(date, Value, .interactive = interactive)

FREQ<- 52
WEEKS <- length(country$Value)

future <- as.integer(WEEKS*0.2)
dlimit <- head(tail(country,future),1)$date

nTrain <- length(country$date) - future

```

```

train <- country %>% select(Value, date) %>% filter(date < dlimit)
valid <- country %>% select(Value, date) %>% filter(date >= dlimit)

# Split Data 80/20
splits <- initial_time_split(train, prop = 0.8)

# Model 1: auto_arima ----
model_fit_arima_no_boost <- arima_reg() %>%
  set_engine(engine = "auto_arima") %>%
  fit(Value ~ date, data = training(splits))

## frequency = 13 observations per 1 quarter

# Model 1b: auto_arima ---- ARIMA(3,1,0)(0,0,2)[52]
model_fit_arima_52 <- arima_reg(seasonal_period = 52,
  non_seasonal_ar = 3, non_seasonal_differences = 1,
  non_seasonal_ma = 0, seasonal_ar = 0, seasonal_differences = 0,
  seasonal_ma = 2) %>%
  set_engine(engine = "arima") %>%
  fit(Value ~ date, data = training(splits))

# Model 2: arima_boost ----
model_fit_arima_boosted <- arima_boost(
  min_n = 1,
  learn_rate = 0.015
) %>%
  set_engine(engine = "auto_arima_xgboost") %>%
  fit(Value ~ date + as.numeric(date) + factor(month(date, label = TRUE), ordered = F),
    data = training(splits))

## frequency = 13 observations per 1 quarter

# Model 3: ets ----
model_fit_ets <- exp_smoothing() %>%
  set_engine(engine = "ets") %>%
  fit(Value ~ date, data = training(splits))

## frequency = 13 observations per 1 quarter

# Model 4: prophet ----
model_fit_prophet <- prophet_reg(seasonality_weekly = TRUE) %>%
  set_engine(engine = "prophet") %>%
  fit(Value ~ date, data = training(splits))

## Disabling daily seasonality. Run prophet with daily.seasonality=TRUE to override this.

# Model 5: lm ----
model_fit_lm <- linear_reg() %>%
  set_engine("lm") %>%
  fit(Value ~ as.numeric(date) + factor(month(date, label = TRUE), ordered = FALSE),
    data = training(splits))

# Model 6: earth ----
model_spec_mars <- mars(mode = "regression") %>% set_engine("earth")

```

```

recipe_spec <- recipe(Value ~ date, data = training(splits)) %>%
  step_date(date, features = "month", ordinal = FALSE) %>%
  step_mutate(date_num = as.numeric(date)) %>%
  step_normalize(date_num) %>%
  step_rm(date)

wflw_fit_mars <- workflow() %>%
  add_recipe(recipe_spec) %>%
  add_model(model_spec_mars) %>%
  fit(training(splits))

```

```
## Loading required package: Formula
```

```
## Loading required package: plotmo
```

```
## Loading required package: plotrix
```

```
##
```

```
## Attaching package: 'plotrix'
```

```
## The following object is masked from 'package:scales':
```

```
##
```

```
##      rescale
```

```
## Loading required package: TeachingDemos
```

```

models_tbl <- modeltime_table(
  model_fit_arima_no_boost,
  model_fit_arima_boosted,
  model_fit_arima_52,
  model_fit_ets,
  model_fit_prophet,
  model_fit_lm,
  wflw_fit_mars
)

```

```
models_tbl
```

```
## # Modeltime Table
```

```
## # A tibble: 7 x 3
```

```
##   .model_id .model      .model_desc
```

```
##         <int> <list>      <chr>
```

```
## 1         1 <fit[+]> ARIMA(5,1,1)(0,0,1)[13]
```

```
## 2         2 <fit[+]> ARIMA(1,1,1)(1,0,0)[13] W/ XGBOOST ERRORS
```

```
## 3         3 <fit[+]> ARIMA(3,1,0)(0,0,2)[52]
```

```
## 4         4 <fit[+]> ETS(M,AD,A)
```

```
## 5         5 <fit[+]> PROPHET
```

```
## 6         6 <fit[+]> LM
```

```
## 7         7 <workflow> EARTH
```

```
# calibrate
```

```
calibration_tbl <- models_tbl %>%
```

```
  modeltime_calibrate(new_data = train)
```

```
calibration_tbl
```


```
## # Modeltime Table
```

```
## # A tibble: 7 x 5
```

```
##   .model_id .model      .model_desc      .type .calibration_data
##       <int> <list>      <chr>           <chr> <list>
## 1         1 <fit[+]>    ARIMA(5,1,1)(0,0,1)[13]    Test <tibble>
## 2         2 <fit[+]>    ARIMA(1,1,1)(1,0,0)[13] W/ XGB00~ Test <tibble>
## 3         3 <fit[+]>    ARIMA(3,1,0)(0,0,2)[52]    Test <tibble>
## 4         4 <fit[+]>    ETS(M,AD,A)             Test <tibble>
## 5         5 <fit[+]>    PROPHET                  Test <tibble>
## 6         6 <fit[+]>    LM                       Test <tibble>
## 7         7 <workflow> EARTH                     Test <tibble>
```

```
# test set forecast and accuracy
calibration_tbl %>%
  modeltime_forecast(
    new_data = testing(splits),
    actual_data = country,
  ) %>%
  plot_modeltime_forecast(
    .legend_max_width = 25, # For mobile screens
    .interactive = interactive,
    .conf_interval_show = FALSE
  )
```

Forecast Plot



Legend

— ACTUAL	— 2_ARIMA(1,1,1)(1,0,0)[...]	— 4_ETS(M,AD,A)	— 6_LM
— 1_ARIMA(5,1,1)(0,0,1)[13]	— 3_ARIMA(3,1,0)(0,0,2)[52]	— 5_PROPHET	— 7_EA

```
calibration_tbl %>%
  modeltime_accuracy() %>%
  table_modeltime_accuracy(
    .interactive = interactive
  )
```


Accuracy Table

.model_id	.model_desc	.type	mae	mape	mase	smape	rmse	r
1	ARIMA(5,1,1)(0,0,1)[13]	Test	6194.18	11.11	7.08	10.68	7718.55	0.0
2	ARIMA(1,1,1)(1,0,0)[13] W/ XGBOOST ERRORS	Test	6185.72	11.10	7.07	10.67	7714.75	0.0
3	ARIMA(3,1,0)(0,0,2)[52]	Test	6695.46	12.10	7.65	11.50	7976.54	0.0
4	ETS(M,AD,A)	Test	6055.22	10.84	6.92	10.45	7617.15	0.0
5	PROPHET	Test	2509.52	3.77	2.87	4.07	5475.89	0.0
6	LM	Test	2695.90	4.08	3.08	4.40	5641.87	0.4
7	EARTH	Test	2358.15	3.57	2.69	3.81	5107.08	0.0


```

refit_tbl <- calibration_tbl %>%
  modeltime_refit(data = train)

## frequency = 13 observations per 1 quarter
## frequency = 13 observations per 1 quarter
## frequency = 13 observations per 1 quarter

## Disabling daily seasonality. Run prophet with daily.seasonality=TRUE to override this.
refit_tbl %>%
  modeltime_forecast(h = "85 weeks", actual_data = country) %>%
  plot_modeltime_forecast(
    .legend_max_width = 25, # For mobile screens
    .interactive       = interactive,
    .conf_interval_show = FALSE
  )

```



Legend

- ACTUAL
- 1_UPDATE: ARIMA(2,1,0)...
- 2_UPDATE: ARIMA(3,1,0)...
- 3_ARIMA(3,1,0)(0,0,2)[52]
- 4_UPDATE: ETS(M,AD,N)
- 5_PROPHET

```
ensemble_fit <- refit_tbl %>%
  ensemble_average(type = "mean")

ensemble_fit

## -- Modeltime Ensemble -----
## Ensemble of 7 Models (MEAN)
##
## # Modeltime Table
## # A tibble: 7 x 5
##   .model_id .model      .model_desc      .type .calibration_data
##   <int> <list>      <chr>      <chr> <list>
## 1       1 <fit[+]> UPDATE: ARIMA(2,1,0)(0,0,1)[13] Test <tibble>
## 2       2 <fit[+]> UPDATE: ARIMA(3,1,0)(1,0,0)[13] ~ Test <tibble>
## 3       3 <fit[+]> ARIMA(3,1,0)(0,0,2)[52] Test <tibble>
## 4       4 <fit[+]> UPDATE: ETS(M,AD,N) Test <tibble>
## 5       5 <fit[+]> PROPHET Test <tibble>
## 6       6 <fit[+]> LM Test <tibble>
## 7       7 <workflow> EARTH Test <tibble>

# Calibration
e.calibration_tbl <- modeltime_table(
  ensemble_fit
) %>%
  modeltime_calibrate(train, quiet = FALSE)

# get split for in-sample forecast
```

```
splits <- initial_time_split(country, prop = 0.8)
# Forecast vs Test Set
#par(mfrow = c(2,1))
```

```
png("docs/ensemble_models.png")
calibration_tbl %>%
  modeltime_forecast(
    new_data = testing(splits),
    actual_data = train,
  ) %>%
  plot_modeltime_forecast(
    .legend_max_width = 25, # For mobile screens
    .interactive = interactive,
    .conf_interval_show = FALSE
  )
dev.off()
```

```
## pdf
## 2
```

```
png("docs/ensemble.png")
e.calibration_tbl %>%
  modeltime_forecast(
    new_data = testing(splits),
    actual_data = country
  ) %>%
  plot_modeltime_forecast(.interactive = FALSE,
    .title = paste0(mtitle, "\n ", stitle),
    .x_lab = x_lab,
    .y_lab = y_lab)
```

```
## Warning in max(ids, na.rm = TRUE): no non-missing arguments to max; returning
## -Inf
```

```
dev.off()
```

```
## pdf
## 2
```

```
#par(mfrow = c(1,1))
```

```
e.calibration_tbl %>%
  modeltime_accuracy() %>%
  table_modeltime_accuracy(
    .interactive = interactive
  )
```

Accuracy Table

.model_id	.model_desc	.type	mae	mape	mase	smape	rmse	rsq
1	ENSEMBLE (MEAN): 7 MODELS	Test	3669.21	6.26	4.19	6.28	5178.64	0.62

```
e.calibration_tbl %>%
  modeltime_accuracy() #>% table()
```

```
## # A tibble: 1 x 9
##   .model_id .model_desc .type mae mape mase smape rmse rsq
##   <int> <chr>          <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1         1 ENSEMBLE (MEAN): 7 MODELS Test 3669. 6.26 4.19 6.28 5179. 0.623
```

Individual Models

Using HoltWinters with trend and seasonality (parameters derived from data)

HoltWinters is the best for the test data (ie last 20%)

```
country.hw<- HoltWinters(train.ts)
country.hw
```

```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = train.ts)
##
## Smoothing parameters:
##   alpha: 0.9376002
##   beta : 0
##   gamma: 1
##
## Coefficients:
##           [,1]
## a    59738.22486
## b      13.09048
## s1  -2821.36446
## s2  -2834.79382
## s3  -3080.29107
## s4  -3224.46196
## s5  -3154.92509
## s6  -3091.05137
## s7  -2785.08032
## s8  -2644.83795
## s9  -2490.27139
## s10 -2228.03700
## s11 -2388.65894
## s12  -872.73337
## s13 -1121.50495
## s14  -734.70776
## s15  -808.60782
## s16  -395.09193
## s17  -320.47990
## s18   238.23181
## s19   213.15821
## s20   532.28066
## s21  2194.28354
## s22  1522.97739
## s23  1673.04283
## s24  1422.59552
## s25  4169.85871
## s26  4190.41023
## s27  3437.99924
```

```
## s28 3417.28944
## s29 3149.68591
## s30 3571.42453
## s31 3761.24021
## s32 3389.39582
## s33 3698.40033
## s34 4171.86394
## s35 2912.11974
## s36 2832.39934
## s37 2432.61027
## s38 2599.34089
## s39 2275.34285
## s40 960.92896
## s41 -486.16957
## s42 -970.60752
## s43 -1278.61882
## s44 -1656.34527
## s45 -1763.14950
## s46 -2423.46344
## s47 -2204.59565
## s48 -2397.96531
## s49 -2564.29923
## s50 -2607.12016
## s51 -2598.92930
## s52 -2519.22486
```

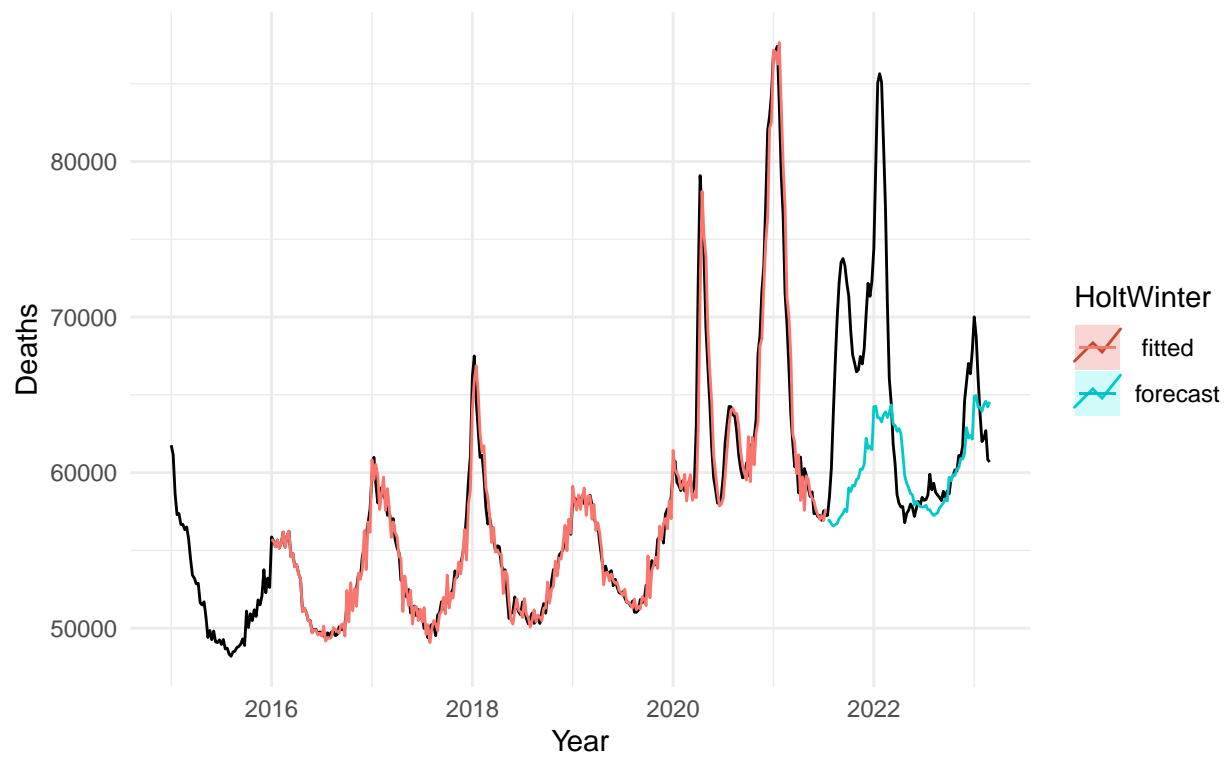
```
country.hw.fc <- forecast(country.hw, h = future)
```

#We see from the plot that the Holt-Winters exponential method is very successful in modelling the seasonality, but the level is off for predictions because not taking enough account of the surge

```
autoplot(country.ts) +
  autolayer(country.hw.fc, PI=F, series="forecast") +
  autolayer(country.hw$fitted[,1], series = "fitted")+
  ggtitle(paste0("Mortality in ", country_code, " HoltWinter Forecast (0.94, 0, 1)",
    subtitle = stitle) +
  theme_minimal() +
  xlab(x_lab) + ylab(y_lab) +
  guides(colour=guide_legend(title="HoltWinter"))
```

Mortality in USA HoltWinter Forecast (0.94, 0, 1)

1/2015 – 2/2023




```
ggsave(here("docs", "holtwinter.png"))
```

```
## Saving 6.5 x 4.5 in image
```

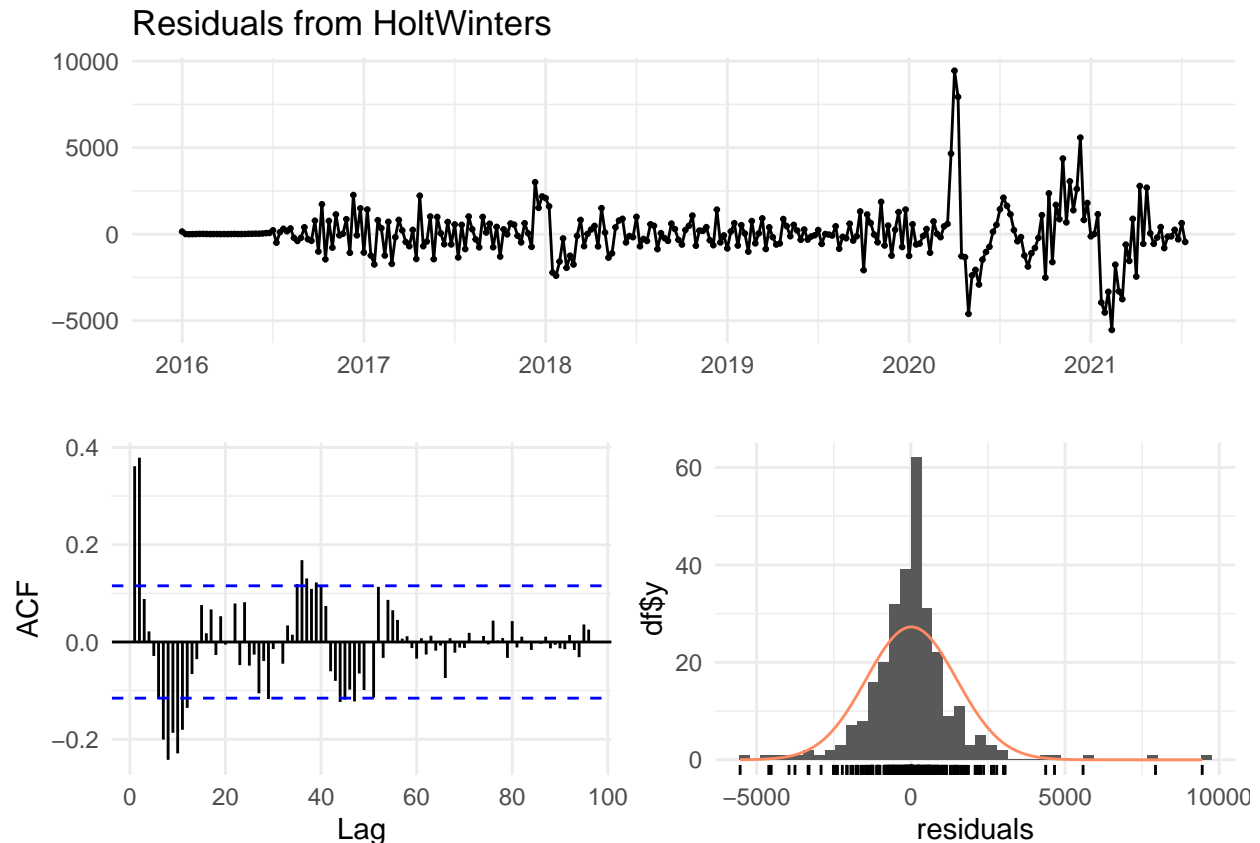
```
Acf(na.omit(country.hw.fc$residuals), lag.max = 20)
```

Series na.omit(country.hw.fc\$residuals)



```
Box.test(country.hw.fc$residuals, lag=20, type = "Ljung-Box")
```

```
##  
## Box-Ljung test  
##  
## data: country.hw.fc$residuals  
## X-squared = 163.76, df = 20, p-value < 2.2e-16  
checkresiduals(country.hw)
```



```
##
##  Ljung-Box test
##
## data:  Residuals from HoltWinters
## Q* = 255.43, df = 58, p-value < 2.2e-16
##
## Model df: 0.   Total lags used: 58
print(forecast::accuracy(country.hw.fc, valid.ts))
```

	ME	RMSE	MAE	MPE	MAPE	MASE
## Training set	16.66519	1465.785	905.4864	0.01273601	1.492413	0.2386424
## Test set	4260.12115	7832.559	5312.6288	5.71928962	7.498161	1.4001520

	ACF1	Theil's U
## Training set	0.3613692	NA
## Test set	0.9494772	4.066744

Linear regression model $\ln(y) \sim \text{trend} + \text{season}$

best result on training, worst on test, so overfitted and over predicts on test, so penalised more by MAPE

<https://www.rdocumentation.org/packages/forecast/versions/8.21/topics/tslm>

```
tslm(train.ts~ season+trend, lambda=NULL) %>% forecast(h=future) %>% forecast::accuracy()
```

	ME	RMSE	MAE	MPE	MAPE	MASE
## Training set	-3.635105e-13	4294.39	3017.928	-0.4148513	5.03524	0.7953799

ACF1


```

## Training set 0.964033
# mape 5.03524

tslm(log(train.ts)~ season+trend, lambda=NULL) %>% forecast(h=future) %>% forecast::accuracy()

##           ME           RMSE           MAE           MPE           MAPE
## Training set 6.270777e-17 0.06448101 0.04687401 -0.003380151 0.4258694
##           MASE           ACF1
## Training set 0.7531913 0.9616223
# mape 0.4258694

# final lm model
country.lm <- tslm(log(train.ts)~ trend + season, lambda=NULL)
country.lm.fc <- forecast(country.lm,h=future)
forecast::accuracy(country.lm.fc, valid.ts)

##           ME           RMSE           MAE           MPE           MAPE
## Training set -4.176698e-17 6.448101e-02 4.687401e-02 -0.003380151 0.4258694
## Test set      6.461939e+04 6.500931e+04 6.461939e+04 99.982671444 99.9826714
##           MASE           ACF1 Theil's U
## Training set 7.531913e-01 0.9616223      NA
## Test set      1.038332e+06 0.9585577 38.23552
# mape: train= 0.4258694, test = 99.9826714

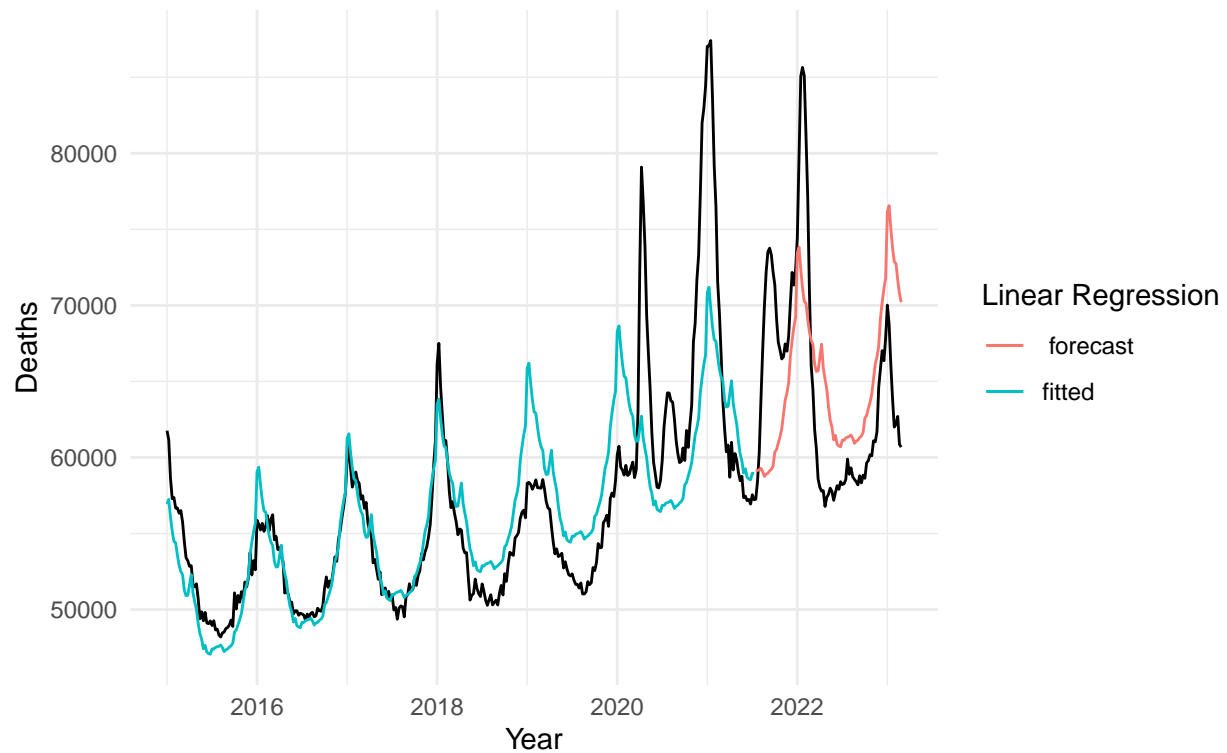
autoplot(country.ts) +
  autolayer(exp(country.lm$fitted.values), PI=F, series="fitted") +
  autolayer(exp(country.lm.fc$mean), series = " forecast")+
  ggtitle(paste0("Mortality in ", country_code, " Linear Regression (ln(y)~s+t)",
    subtitle = stitle) +
  theme_minimal() +
  xlab(x_lab) + ylab(y_lab) +
  guides(colour=guide_legend(title="Linear Regression"))

## Warning in ggplot2::geom_line(ggplot2::aes(x = .data[["timeVal"]], y =
## .data[["seriesVal"]], : Ignoring unknown parameters: `PI`

```

Mortality in USA Linear Regression ($\ln(y) \sim s+t$)

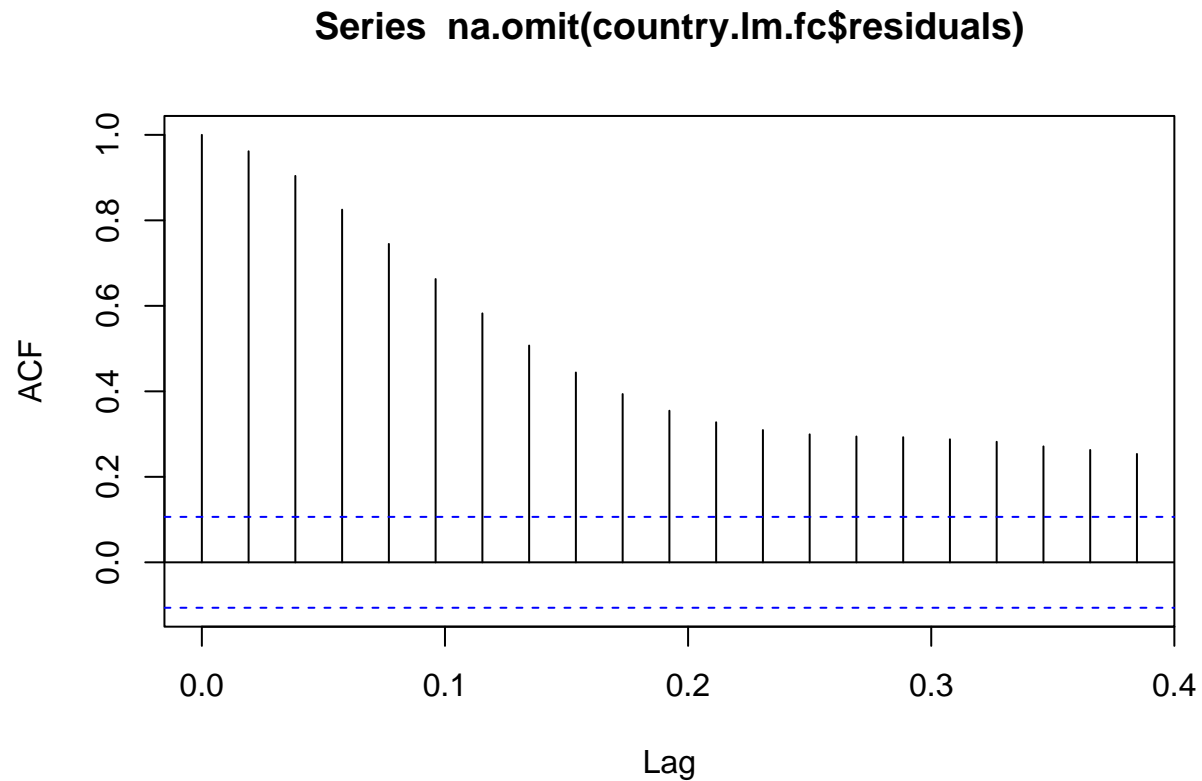
1/2015 – 2/2023



```
country.lm
```


```
##
## Call:
## tslm(formula = log(train.ts) ~ trend + season, lambda = NULL)
##
## Coefficients:
## (Intercept)      trend    season2    season3    season4    season5
##  10.949092    0.000699    0.004547   -0.016308   -0.033887   -0.047049
##    season6    season7    season8    season9    season10    season11
##  -0.049354   -0.065235   -0.078424   -0.086962   -0.091705   -0.112080
##    season12    season13    season14    season15    season16    season17
##  -0.119692   -0.119851   -0.107349   -0.094909   -0.118964   -0.130976
##    season18    season19    season20    season21    season22    season23
##  -0.141607   -0.160590   -0.175472   -0.183548   -0.197699   -0.193596
##    season24    season25    season26    season27    season28    season29
##  -0.203943   -0.206333   -0.208037   -0.201393   -0.202446   -0.200515
##    season30    season31    season32    season33    season34    season35
##  -0.200475   -0.200368   -0.199401   -0.203820   -0.209925   -0.208137
##    season36    season37    season38    season39    season40    season41
##  -0.207721   -0.205808   -0.204761   -0.201643   -0.187132   -0.185217
##    season42    season43    season44    season45    season46    season47
##  -0.178645   -0.172756   -0.164975   -0.150572   -0.136308   -0.129738
##    season48    season49    season50    season51    season52
##  -0.119964   -0.093855   -0.079548   -0.067418   -0.058610
```

```
acf(na.omit(country.lm.fc$residuals), lag.max = 20)
```



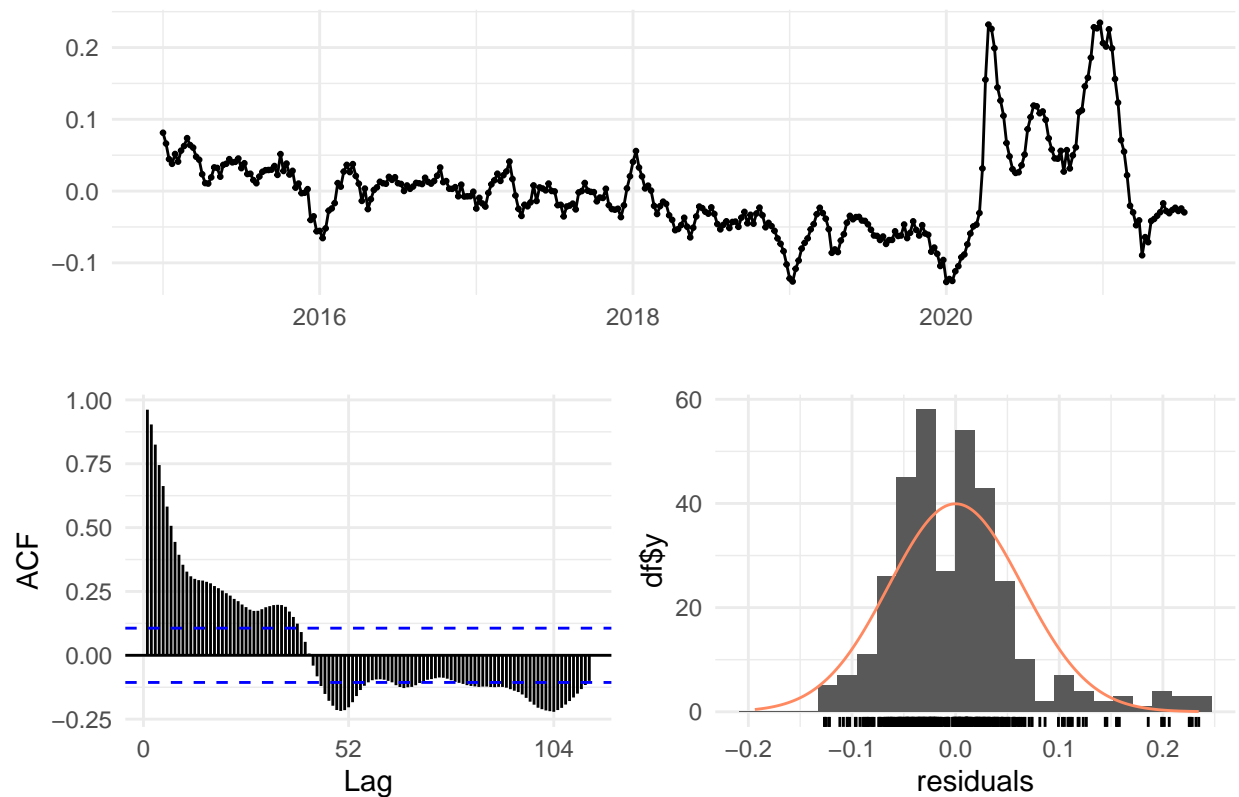
```
Box.test(country.lm.fc$residuals, lag=20, type = "Ljung-Box")
```

```
##  
## Box-Ljung test  
##  
## data: country.lm.fc$residuals  
## X-squared = 1851.9, df = 20, p-value < 2.2e-16  
ggtsdisplay(country.lm.fc$residuals)
```



```
checkresiduals(country.lm)
```

Residuals from Linear regression model



```
##
## Breusch-Godfrey test for serial correlation of order up to 68
##
## data: Residuals from Linear regression model
## LM test = 323.72, df = 68, p-value < 2.2e-16

# can't use default train.ts because returns error for frequency=52
ets.ts <- ts(country$Value, start=c(startYear,startWeek), frequency=13)
ets.train.ts <- window(ets.ts, start = c(startYear, startWeek),
                      end = c(startYear, nTrain))
ets.valid.ts <- window(ets.ts, start = c(startYear, nTrain+1),
                     end = c(startYear, nTrain+nValid))

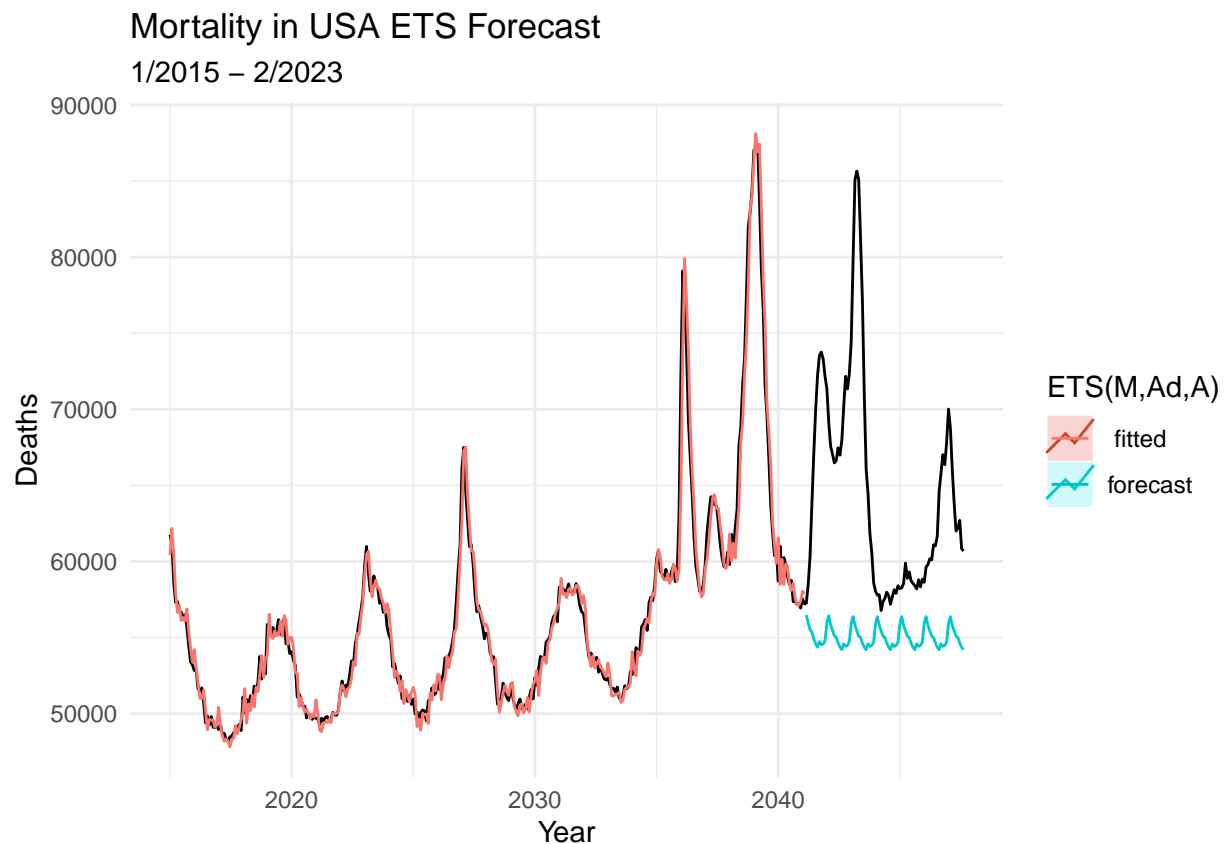
# if use ZZZ don't get any seasonality
a.ets <- ets(ets.train.ts, model="ZZZ", alpha = NULL)
a.ets <- ets(ets.train.ts, model="ZZA", alpha = NULL)

# get prediction
a.ets.fc <- forecast(a.ets, h = future, level = 0)
a.ets

## ETS(M,Ad,A)
##
## Call:
## ets(y = ets.train.ts, model = "ZZA", alpha = NULL)
##
## Smoothing parameters:
```

```
##      alpha = 0.9997
##      beta  = 0.124
##      gamma = 1e-04
##      phi   = 0.8001
##
##      Initial states:
##      l = 59748.0635
##      b = -286.6304
##      s = -301.2254 -524.9653 -585.354 -413.9644 -801.7451 -633.9307
##           -363.1328 5.6935 94.7059 437.1969 804.6165 1371.303 910.8021
##
##      sigma: 0.0197
##
##      AIC      AICc      BIC
## 6764.042 6766.417 6836.792
```

```
autoplot(ets.ts) +
  autolayer(a.ets.fc, PI=T, series="forecast") +
  autolayer(a.ets$fitted, series = "fitted")+
  ggtitle(paste0("Mortality in ", country_code, " ETS Forecast "),
    subtitle = stitle) +
  theme_minimal() +
  xlab(x_lab) + ylab(y_lab) +
  guides(colour=guide_legend(title="ETS(M,Ad,A)"))
```



```
forecast::accuracy(a.ets.fc, ets.valid.ts)
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  -7.072029 1205.222  815.8181 -0.01508594  1.385793 0.1534662
## Test set      9592.288290 11946.367 9592.2883 13.90043923 13.900439 1.8044361
##           ACF1 Theil's U
## Training set 0.3288586      NA
## Test set     0.9558169  6.330288
```

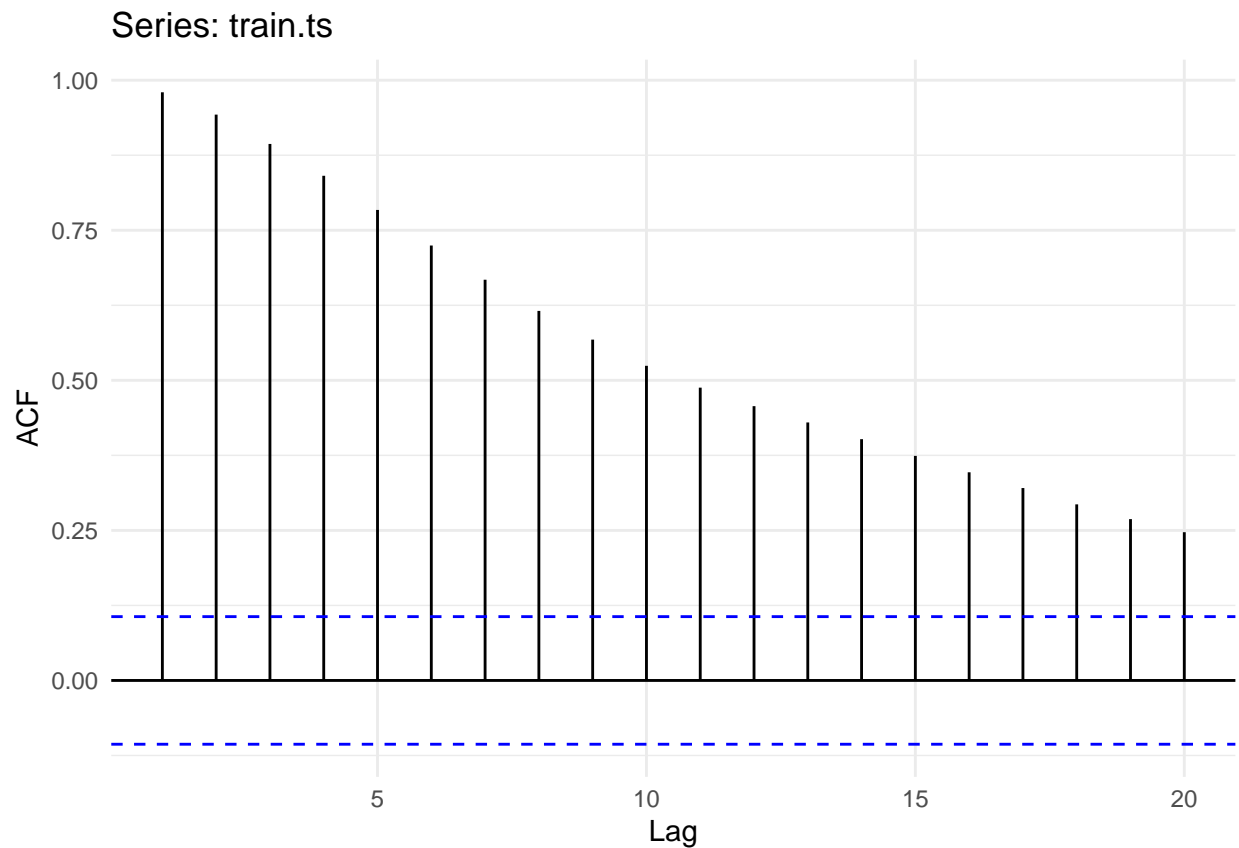
Autocorrelation

The ACF values are all above the threshold. The ACF of the time series is decreasing slowly, which suggests that it is non-stationary.


You can't predict values for non-stationary data.

There is a clear seasonal (and cyclical) effect, (Shmueli and Lichtendahl 2016)

```
# look at autocorrelation
ggAcf(train.ts, lag=20)
```



```
ggtsdisplay(train.ts)
```



ARIMA

An AR (autoregressive) model is usually used to model a time series which shows longer term dependencies between successive observations. Intuitively, it makes sense that an AR model could be used to describe the time series of mortality, as we would expect some factors which affect mortality rates in one year to affect those in later years.

```
tunes <- readRDS(file="data/arimaManualFit.rds")
```

```
best <- tunes[tunes$mape==min(tunes$mape),1:6] %>% as.double()
print(best)
```

```
## [1] 1 0 2 2 2 2
```

```
tunes %>% arrange(mape) %>% head()
```

```
## # A tibble: 6 x 15
##       a     b     c     d     e     f     aic     aicc     me     rmse     mae     mpe
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     1     0     2     2     2     2  5714.  5715.   38.5  873.   555.  0.0482
## 2     2     0     1     2     2     2  5725.  5725.   57.0  897.   562.  0.0755
## 3     1     0     1     2     2     2  5767.  5767.   40.5  937.   587.  0.0500
## 4     1     0     0     2     2     2  5821.  5822.   39.4 1034.   618.  0.0473
## 5     0     1     2     2     2     2  5710.  5710.  -19.9  968.   606. -0.0308
## 6     0     1     1     2     2     2  5758.  5758. -30.8 1007.   621. -0.0546
## # i 3 more variables: mape <dbl>, mase <dbl>, acf1 <dbl>
```



```

# fit best manual arima model
#MANUAL
country.am.fc <- (country.am <- arima(train.ts, order = c(best[1:3]),
                                     seasonal = c(best[4:6]))) %>% forecast()

##AUTO
country.aa.fc <-(country.aa <- auto.arima(train.ts)) %>% forecast()
#ARIMA(3,1,0)(0,0,2)[52]

# accuracy scores
print(forecast::accuracy(country.am.fc, valid.ts))

##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  49.28524 1013.021  626.5231  0.06686918  1.058339  0.1651212
## Test set     -791.75378 8580.879 6670.1722 -1.97661934 10.051741  1.7579348
##           ACF1 Theil's U
## Training set  0.03455536      NA
## Test set     0.96428426  4.881296

print(forecast::accuracy(country.aa.fc, valid.ts))


##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -5.377429 1160.430  791.2118 -0.01086611  1.357587  0.2085252
## Test set     6457.400499 9169.264 6647.6901  9.10709301  9.438398  1.7520096
##           ACF1 Theil's U
## Training set  0.003307442      NA
## Test set     0.955775638  4.761724

#https://www.educba.com/arima-model-in-r/

```

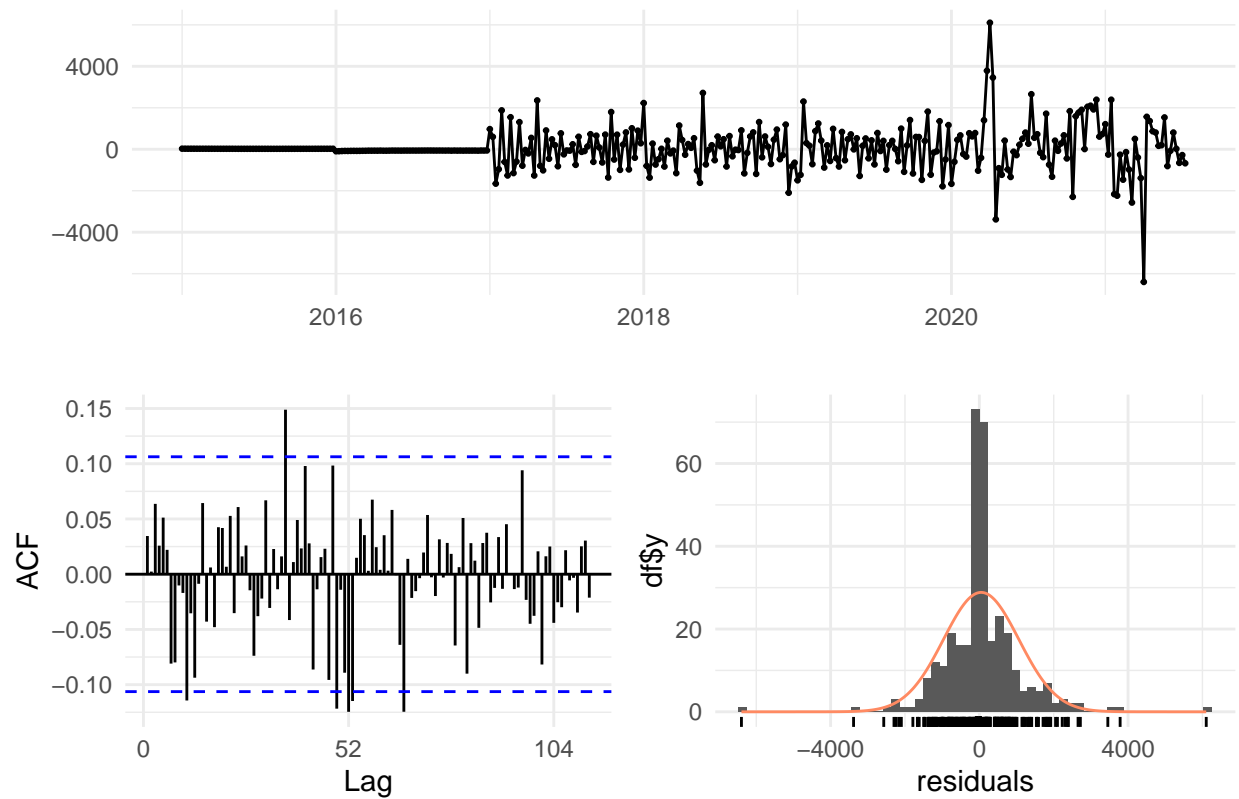
Plotting the arima model:

```
ggtsdisplay(resid(country.am))
```




```
checkresiduals(country.am)
```


Residuals from ARIMA(1,0,2)(2,2,2)[52]



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,0,2)(2,2,2)[52]
## Q* = 88.947, df = 61, p-value = 0.01128
##
## Model df: 7.   Total lags used: 68
ggtsdisplay(resid(country.aa))
```



```
checkresiduals(country.aa)
```




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(3,1,0)(0,0,2)[52]
## Q* = 89.649, df = 63, p-value = 0.01535
##
## Model df: 5.   Total lags used: 68

#-----
# compare forecasts for individual models
autoplot(country.ts) +
  autolayer(country.aa.fc, PI=F, series="Auto") +
  autolayer(country.am.fc, PI=F, series = "Manual")+
  autolayer(exp(country.lm.fc$mean), PI=F, series = "log(y)~t+s")+
  autolayer(country.hw.fc, PI=F, series = "HoltWinters")+
  ggtitle("Comparison of forecasts",
    subtitle = stitle) +
  theme_minimal() +
  xlab(x_lab) + ylab(y_lab) +
  guides(colour=guide_legend(title="Model"))

## Warning in ggplot2::geom_line(ggplot2::aes(x = .data[["timeVal"]], y =
## .data[["seriesVal"]], : Ignoring unknown parameters: `PI`
```

Comparison of forecasts 1/2015 – 2/2023




Appendix 1 - training on pre-COVID data only


```
covid <- length(country[country$YEAR>= 2020,]$date)

nTrain <- length(country.ts) - covid
train.ts <- window(country.ts, start = c(startYear, startWeek),
                    end = c(startYear, nTrain))
covid.ts <- window(country.ts, start = c(startYear, nTrain+1),
                   end = c(startYear, nTrain+covid))

# restricting the decomposition to pre-covid changes the shape of the seasonal component slightly and m
covid.ts %>% stl(s.window="periodic") %>% autoplot()
```



```
train.ts %>% stl(s.window="periodic") %>% autoplot()
```



```
#summary(fit.lm)
# adj r^2 = 0.8866
# significant trend coefficient 2.719e-04
#fit.lm$coefficients[2]

covid.arima <- auto.arima(train.ts) %>% forecast(covid)
covid.hw <- HoltWinters(train.ts) %>% forecast(covid)
covid.lm <- tslm(log(train.ts)~ trend + season, lambda=NULL) %>% forecast(covid)
```

```
## Warning in forecast.lm(., covid): newdata column names not specified,
## defaulting to first variable required.
```

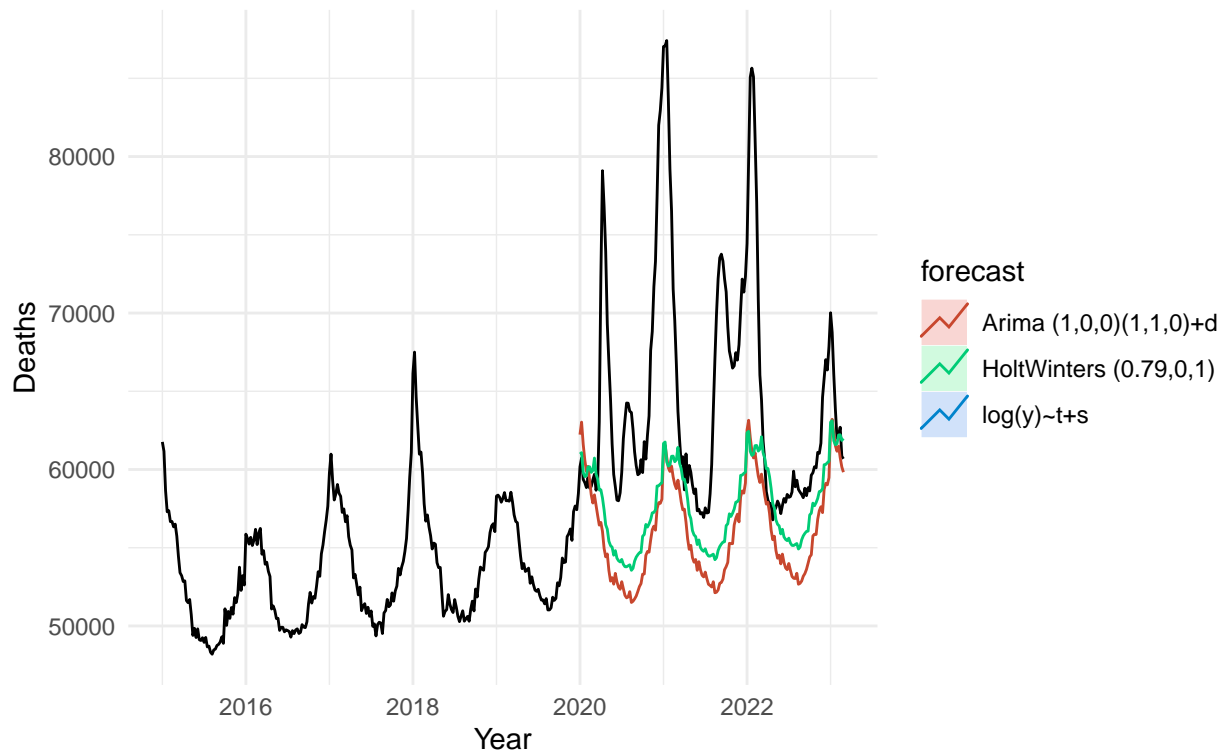
```
# plot model forecasts
autoplot(country.ts) +
  autolayer(covid.arima, PI=F, series="Arima (1,0,0)(1,1,0)+d") +
  autolayer(exp(covid.lm$mean), PI=F, series = "log(y)~t+s")+
  autolayer(covid.hw, PI=F, series = "HoltWinters (0.79,0,1)") +
  ggtitle("Comparison of forecasts",
    subtitle = stitle) +
  theme_minimal() +
  xlab(x_lab) + ylab(y_lab) +
  guides(colour=guide_legend(title="forecast"))
```

```
## Warning in ggplot2::geom_line(ggplot2::aes(x = .data[["timeVal"]], y =
## .data[["seriesVal"]], : Ignoring unknown parameters: `PI`

## `geom_line()`: Each group consists of only one observation.
## i Do you need to adjust the group aesthetic?
```


Comparison of forecasts

1/2015 – 2/2023



```
ggsave(here("docs", "noCovid.png"))
```

```
## Saving 6.5 x 4.5 in image
## `geom_line()`: Each group consists of only one observation.
## i Do you need to adjust the group aesthetic?
```

All the models fit well to the training data, and are consistent in predicting the mortality that would have been expected in 2020+ if COVID-19 hadn't happened, so they could be used to estimate the excess mortality in the USA due to COVID.

Thu May 4 00:12:59 2023

refs:

<https://github.com/FinYang/tsdl/tree/master> The Time Series Data Library (TSDL) was created by Rob Hyndman, Professor of Statistics at Monash University, Australia.

<https://www.rdocumentation.org/packages/forecast/versions/8.21/topics/tslm>

Coghlan, Avril. 2023. "Welcome to a Little Book of r for Time Series!" 2023. <https://a-little-book-of-r-for-time-series.readthedocs.io/en/latest/>.

Dancho, Matt. 2023. *Modeltime: The Tidymodels Extension for Time Series Modeling*.

Hyndman, Rob J, and George Athanasopoulos. 2023. *Forecasting: Principles and Practice*. 3rd ed. Monash University. <https://otexts.com/fpp3/>.

OECD. 2023. "COVID-19 Health Indicators, Mortality (by Week)." <https://doi.org/https://doi.org/https://doi.org/10.1787/cd2bda32-en>.

Shmueli, Galit, and Kenneth C Lichtendahl. 2016. *Practical Time Series Forecasting with r: A Hands-on Guide [2nd Edition]*. Practical Analytics. Axelrod Schnall Publishers. <https://books.google.ie/books?id=mxWXdWAAQBAJ>.